

Current

A PM'S OPERATING SYSTEM

From IC to Orchestrator

How I run product with a team of AI agents

Albert Collell · Product, Current

THE STARTING POINT

One PM, a whole product function

No researcher, no analyst, no designer, no PMM. Just the jobs.

Research synthesis

Dozens of calls a month, and someone has to turn them into insight

Usage analytics

Real numbers behind every outcome claim

Prototyping and design

Something clickable before engineering commits

Engineering feasibility

A sparring partner on the how, before a brief lands

GTM and positioning

Battlecards, release notes, messaging

Operations

Meeting notes, action items, follow-ups, recaps

THE SHIFT

From doing the work to directing it

BEFORE · AS AN IC


- I draft every PRD
- I write up every meeting
- I chase actions from memory
- Competitive research, when I can
- My judgment spread thin

NOW · AS AN ORCHESTRATOR

- Agents draft, I review outcomes
- Meetings synthesize themselves
- 472 actions tracked in one place
- Competitive sweep on a schedule
- Judgment where only I add value

THE SYSTEM


A repo as the company's second brain



Plain markdown in git, run from Claude Code. **Every agent reads the same context before it acts.**

THE WHOLE PICTURE

Business intelligence, end to end



THE TEAM

Nine agents, each standing in for a role

discovery-researcher

Research IC: turns transcripts into opportunity assessments

data-analyst

Analytics: real PostHog numbers, or an honest not instrumented

prototype-engineer

Frontend IC: clickable prototypes in real code

apple-design-reviewer

Design critique against Apple's interface guidelines

solutions-architect

Engineering counterpart: options and feasibility, after the brief

gtm-writer

PMM: battlecards, positioning, release notes. Drafts only

editor

Distills recurring ideas into lessons and public drafts

product-coach


Pressure-tests my docs: Cagan, Hoffman, Ries

cpo-coach

Pressure-tests me, using evidence from the repo

THE OPERATING LOOP

Meetings become owned state within a day



CONTEXT, NOT CONTROL

Explicit lines on who decides what

Agents do alone	Agents flag to me	Only I do
File action items to the tracker	Anything decision-worthy	Log a decision and its rationale
Add customer signal to insights	Customer or stakeholder changes	Edit ground truth in context/
Draft briefs, recaps, battlecards	Cross-area trade-offs	Send email, Slack, or invites
Run small bets within scope	Customer commitments and SLAs	Set OKR targets each quarter

Written into CLAUDE.md and each AGENT.md, so no agent has to guess.

FOUR MONTHS IN · JUNE TO SEPTEMBER 2026

What the system has produced

175

meetings synced from Granola

191

research notes filed by topic

472

action items tracked with owners

11

opportunity assessments

12

decisions logged with rationale

9

specialist agents, 12 commands

That's output, not outcome. The real test is where my hours go.

THE COACHING LAYER

Agents that push back, not just produce

PRODUCT-COACH**Reviews the decision on the page**

- Discovery before delivery (Cagan)
- Speed as strategy (Hoffman)
- Validated learning (Ries)
- Socratic by default, verdict on request

CPO-COACH**Reviews the person making it**

- Grounded in my decisions, tracker, 1:1s
- Where am I not delegating?
- Portfolio calls, managing up and sideways
- Calls product-coach when it's the doc

I also run them on peers' docs, to help them improve, not to grade them.

WHAT I LEARNED

Four lessons from building it

1 Context is the product

CLAUDE.md is a job description. Vague context gives you vague agents.

2 Synthesis is the bottleneck

Capture was solved on day one. Value only showed up once an agent owned the loop.

3 Draw the lines first

File versus flag decided how much I could trust it, and how much I could let go.

4 My job moved to judgment

Approve, kill, set context, be in the room. Everything else is a first draft.

THE IDEA UNDERNEATH

The moat is the system, not the model.

Everyone has the same models. What compounds is the context, the loop, and the lines around them.

TRY IT AT HOME · 30 MINUTES

Build your first loop tonight

1 Open an empty folder

Start Claude Code there. 5 minutes.

2 Paste the prompt

Answer its interview honestly. 15 minutes.

3 Feed it one real meeting

Drop in a transcript, run /synthesize. 10 minutes.

Bonus: ask it for a coach agent that pushes back on your next doc.

PASTE THIS INTO CLAUDE CODE

```
I'm a product manager. Help me build a PM copilot in this folder, one step at a time. First, interview me: my company, product, users, stakeholders, and where my week goes. Then create a CLAUDE.md that says who I am and what to read first; context/ files from my answers; research/ for meeting notes; actions/tracker.md; and a /synthesize command that turns a meeting note into action items (owner and due date) and customer insights. Write down what you may do alone and what you must ask me first. Check with me before each part.
```

START TOMORROW

Three files to begin with

1

Write your CLAUDE.md

Who you are, what matters,
what to read first.

2

Automate one loop you hate

Meeting notes into action items
is a good first one.

3

Write one authority rule

What the agent may do alone,
and what it must bring to you.

Thank you. Questions?

Starter repo: github.com/albertcollell-current/pm-copilot-starter